

## Machine Learning Lab Manual

### BCA 6<sup>th</sup> Sem

LIST OF PROGRAMS
1. Install and set up Python and essential libraries like NumPy and pandas
2. Introduce scikit-learn as a machine learning library. 3. Install and set up scikit-learn and other necessary tools.
4. Write a program to Load and explore the dataset of .CSV and excel files using pandas.
5. Write a program to Visualize the dataset to gain insights using Matplotlib or Seaborn by plotting scatter plots, bar charts.
6. Write a program to Handle missing data, encode categorical variables, and perform feature scaling.
7. Write a program to implement a k-Nearest Neighbours (k-NN) classifier using scikit-learn and Train the classifier on the dataset and evaluate its performance.
8. Write a program to implement a linear regression model for regression tasks and Train the model on a dataset with continuous target variables.
9. Write a program to implement a decision tree classifier using scikit-learn and visualize the decision tree and understand its splits.
10. Write a program to Implement K-Means clustering and Visualize clusters.

## 1. install and set up Python and essential libraries like NumPy and Pandas.

Install Python: If you haven't already installed Python, you can download it from the official website:

To verify (terminal)

```
python --version
```

Install pip: pip is a package manager for Python that allows you to easily install and manage libraries. Most recent versions of Python come with pip pre-installed. You can verify if pip is installed by running the following command in your terminal or command prompt:

```
pip --version
```

Install NumPy and pandas: Once you have Python and pip installed, you can use pip to install NumPy and pandas by running the following commands in your terminal or command prompt:

**#In terminal**

```
pip install numpy
```

```
pip install pandas
```

This will download and install NumPy and Pandas along with any dependencies they require. Verify installation: After installing NumPy and pandas, you can verify that they were installed correctly by running the following commands in Python's interactive mode or a Python script:

```
import numpy
```

```
import pandas
```

```
print(numpy.__version__)
```

```
print(pandas.__version__)
```

These commands should print the versions of NumPy and pandas that were installed.

## Output:

```
(/opt/anaconda3) apple@Apples-MacBook-Pro ~ % python --version
Python 3.9.7
(/opt/anaconda3) apple@Apples-MacBook-Pro ~ % pip --version
pip 21.2.4 from /opt/anaconda3/lib/python3.9/site-packages/pip (python 3.9)
(/opt/anaconda3) apple@Apples-MacBook-Pro ~ % pip install numpy

Requirement already satisfied: numpy in /opt/anaconda3/lib/python3.9/site-packages (1.22.4)
(/opt/anaconda3) apple@Apples-MacBook-Pro ~ % pip install pandas
Requirement already satisfied: pandas in /opt/anaconda3/lib/python3.9/site-packages (1.3.4)
Requirement already satisfied: python-dateutil>=2.7.3 in /opt/anaconda3/lib/python3.9/site-packages (from pandas) (2.8.2)
Requirement already satisfied: pytz>=2017.3 in /opt/anaconda3/lib/python3.9/site-packages (from pandas) (2021.3)
Requirement already satisfied: numpy>=1.17.3 in /opt/anaconda3/lib/python3.9/site-packages (from pandas) (1.22.4)
Requirement already satisfied: six>=1.5 in /opt/anaconda3/lib/python3.9/site-packages (from python-dateutil>=2.7.3->pandas) (1.16.0)
(/opt/anaconda3) apple@Apples-MacBook-Pro ~ %
```

```
|: import numpy
import pandas

print(numpy.__version__)
print(pandas.__version__)
```

1.22.4

1.3.4

## 2. Introduce sci-kit-learn as a machine learning library.

Scikit-learn (Sklearn) is the most useful and robust library for machine learning in Python. It provides a selection of efficient tools for machine learning and statistical modeling including classification, regression, clustering and dimensionality reduction via a consistency interface in Python. This library, which is largely written in Python, is built upon **NumPy**, **SciPy** and **Matplotlib**.

### Installation

If you already installed NumPy and Scipy, the following are the two easiest ways to install scikit-learn –

#### Using pip

The following command can be used to install sci-kit-learn via pip

**pip install -U scikit-learn**

### Features

Rather than focusing on loading, manipulating and summarising data, Scikit-learn library is focused on modeling the data. Some of the most popular groups of models provided by Sklearn are as follows –

**Supervised Learning algorithms** – Almost all the popular supervised learning algorithms, like Linear Regression, Support Vector Machine (SVM), Decision Tree etc., are the part of scikit-learn.

**Unsupervised Learning algorithms** – On the other hand, it also has all the popular unsupervised learning algorithms from clustering, factor analysis, PCA (Principal Component Analysis) to unsupervised neural networks.

**Clustering** – This model is used for grouping unlabeled data.

**Cross Validation** – It is used to check the accuracy of supervised models on unseen data.

### 3. Install and set up scikit-learn and other necessary tools.

scikit-learn, a powerful Python library for machine learning. Here are the steps to set it up:  
Install Python: If you haven't already installed Python, download and install the latest version of Python 3 from the [official Python website](https://www.python.org/).

Install scikit-learn using pip: Open your terminal or command prompt and run the following command:

```
pip install -U scikit-learn
```

To verify your installation, you can use the following commands:

```
python -m pip show scikit-learn
```

# To see which version and where scikit-learn is installed

```
python -m pip freeze
```

# To see all packages installed in the active virtual environment

```
import sklearn
```

```
import numpy
```

```
import pandas
```

```
import matplotlib
```

```
print(sklearn.__version__)
```

```
print(numpy.__version__)
```

```
print(pandas.__version__)
```

```
print(matplotlib.__version__)
```

These commands should print the versions of scikit-learn and other libraries that were installed.

## Output:

```
|: import sklearn
import numpy
import pandas
import matplotlib
print(sklearn.__version__)
print(numpy.__version__)
print(pandas.__version__)
print(matplotlib.__version__)
```

1.0.2

1.22.4

1.3.4

3.4.3

4. Write a program to Load and explore the dataset of .CSV and excel files using pandas.

```
import pandas as pd

def explore_dataset(file_path):

    # Check if the file is a CSV or Excel file

    if file_path.endswith('.csv'):

        # Load CSV file into a pandas DataFrame

        df = pd.read_csv(file_path)

    elif file_path.endswith('.xlsx'):

        # Load Excel file into a pandas DataFrame

        df = pd.read_excel(file_path)

    else:

        print("Unsupported file format. Please provide a CSV or Excel file.")

        return

    # Display basic information about the DataFrame

    print("Dataset information:")

    print(df.info())

    # Display the first few rows of the DataFrame

    print("\nFirst few rows of the dataset:")

    print(df.head())
```

```
# Display summary statistics for numerical columns
```

```
print("\nSummary statistics:")
```

```
print(df.describe())
```

```
# Display unique values for categorical columns
```

```
print("\nUnique values for categorical columns:")
```

```
for column in df.select_dtypes(include='object').columns:
```

```
    print(f"{column}: {df[column].unique()}")
```

```
# Example usage
```

```
file_path = 'IRIS.csv'
```

```
# Change this to the path of your CSV or Excel file
```

```
explore_dataset(file_path)
```



## Output:

---

```
Dataset information:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   sepal_length    150 non-null    float64
1   sepal_width     150 non-null    float64
2   petal_length    150 non-null    float64
3   petal_width     150 non-null    float64
4   species         150 non-null    object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
None
```

First few rows of the dataset:

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

Summary statistics:

	sepal_length	sepal_width	petal_length	petal_width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.054000	3.758667	1.198667
std	0.828066	0.433594	1.764420	0.763161
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

Unique values for categorical columns:

```
species: ['Iris-setosa' 'Iris-versicolor' 'Iris-virginica']
```

**5. Write a program to Visualize the dataset to gain insights using Matplotlib or Seaborn by plotting scatter plots, and bar charts.**

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

def visualize_dataset(file_path):

    # Load the dataset into a pandas DataFrame
    df = pd.read_csv(file_path)

    # Assuming it's a CSV file, change accordingly if it's an Excel file

    # Plot scatter plots

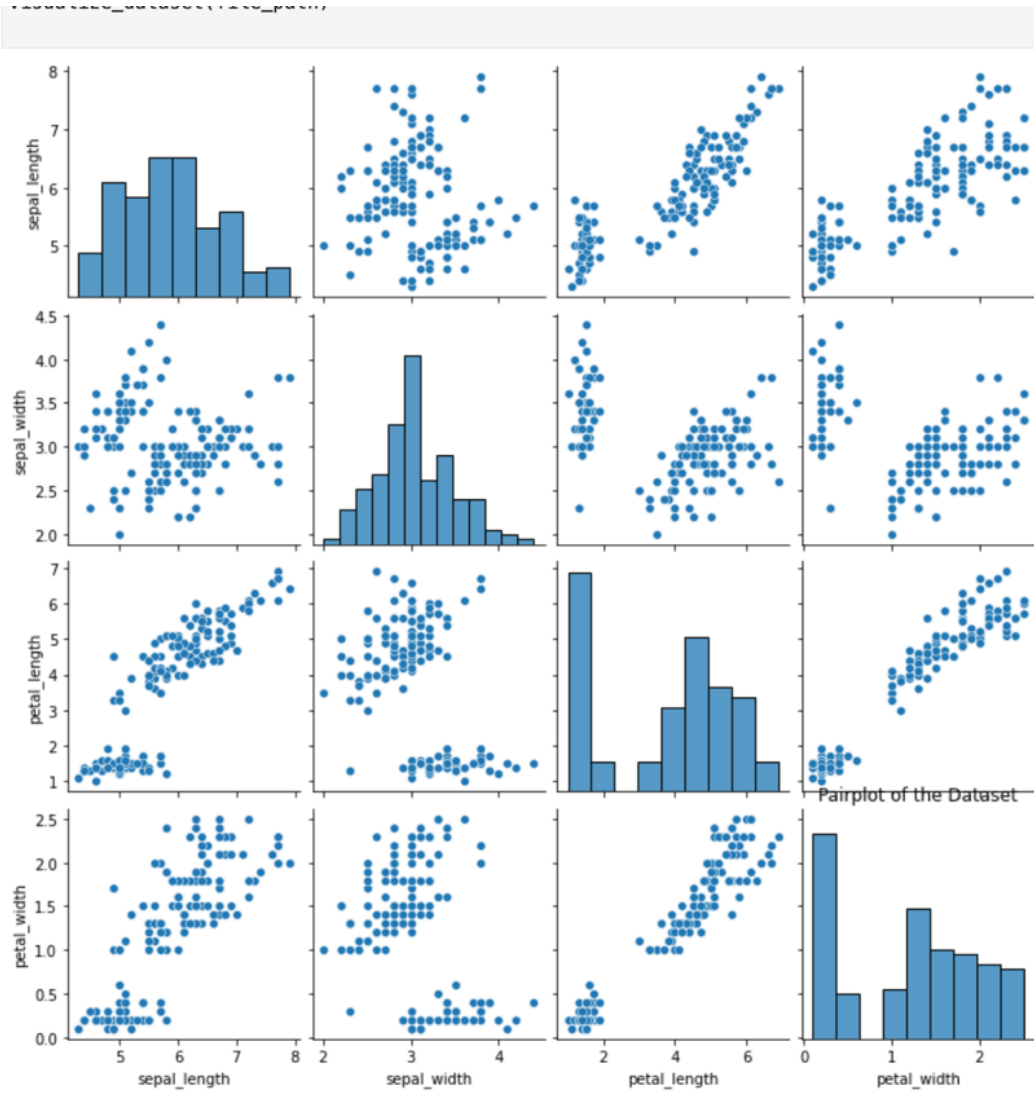
    sns.pairplot(df)
    plt.title("Pairplot of the Dataset")
    plt.show()

    # Plot bar chart for categorical column (assuming the first column is categorical)

    if df.iloc[:, 0].dtype == 'object':
        sns.countplot(x=df.columns[0], data=df)
        plt.title("Bar Chart of Categorical Column")
        plt.xlabel(df.columns[0])
        plt.ylabel("Count")
        plt.show()
    else:
        print("No categorical column found to plot bar chart.")

# Example usage
file_path = 'IRIS.csv' # Change this to the path of your CSV file
visualize_dataset(file_path)
```

## Output:



No categorical column found to plot bar chart.

6. Write a program to Handle missing data, encode categorical variables, and perform feature scaling.

```
import pandas as pd

from sklearn.datasets import load_iris

from sklearn.impute import SimpleImputer

from sklearn.preprocessing import OneHotEncoder, StandardScaler


# Load Iris dataset

iris = load_iris()

iris_df = pd.DataFrame(data=iris.data, columns=iris.feature_names)

iris_df['target'] = iris.target


def preprocess_dataset(df):

    # Handle missing data (Iris dataset doesn't have missing values, but we'll simulate
    some)

    df.iloc[:, 0] = float('NaN')

    # Simulate missing values in the first column

    imputer = SimpleImputer(strategy='mean')

    df[df.columns] = imputer.fit_transform(df[df.columns])

    # Encode categorical variable (if applicable)
```

```
# Since Iris dataset doesn't have categorical variables, we'll skip this step
```

```
# Perform feature scaling
```

```
scaler = StandardScaler()
```

```
df[df.columns[:-1]] = scaler.fit_transform(df[df.columns[:-1]])
```

```
return df
```

```
# Preprocess Iris dataset
```

```
preprocessed_df = preprocess_dataset(iris_df)
```

```
# Display preprocessed dataset
```

```
print("Preprocessed dataset:")
```

```
print(preprocessed_df.head())
```

**Output:**

```
Preprocessed dataset:
  sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)  \
0      1.132764e-15      1.019004      -1.340227      -1.315444
1     -1.196022e+00     -0.131979      -1.340227      -1.315444
2     -1.451098e+00      0.328414      -1.397064      -1.315444
3     -1.578636e+00      0.098217      -1.283389      -1.315444
4     -1.068484e+00      1.249201      -1.340227      -1.315444

  target
0      0.0
1      0.0
2      0.0
3      0.0
4      0.0
```

[ ]:

7. **Write a program to implement a k-Nearest Neighbours (k-NN) classifier using scikit-learn and Train the classifier on the dataset and evaluate its performance.**

```
import numpy as np

import pandas as pd

from sklearn.datasets import load_iris

from sklearn.model_selection import train_test_split

from sklearn.neighbors import KNeighborsClassifier

from sklearn.metrics import accuracy_score, classification_report


# Load Iris dataset

iris = load_iris()

X = iris.data

y = iris.target


# Split the dataset into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)


# Initialize the k-NN classifier

k = 3 # Number of neighbors
```

```
knn_classifier = KNeighborsClassifier(n_neighbors=k)
```

```
# Train the classifier
```

```
knn_classifier.fit(X_train, y_train)
```

```
# Make predictions on the testing set
```

```
y_pred = knn_classifier.predict(X_test)
```

```
# Evaluate the classifier's performance
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
print("Accuracy:", accuracy)
```

```
# Display classification report
```

```
print("Classification Report:")
```

```
print(classification_report(y_test, y_pred, target_names=iris.target_names))
```



## Output:

Accuracy: 1.0

Classification Report:

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	10
versicolor	1.00	1.00	1.00	9
virginica	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

8. **Write a program to implement a linear regression model for regression tasks and Train the model on a dataset with continuous target variables.**

```
import numpy as np

import pandas as pd

from sklearn.datasets import load_boston

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression

from sklearn.metrics import mean_squared_error, r2_score

# Load Boston Housing dataset

boston = load_boston()

X = boston.data

y = boston.target

# Convert the data to a pandas DataFrame for easier manipulation

boston_df = pd.DataFrame(data=X, columns=boston.feature_names)

boston_df['target'] = y

# Split the dataset into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

# Initialize Linear Regression model

linear_regression = LinearRegression()

# Train the model
```

```
linear_regression.fit(X_train, y_train)

# Make predictions on the testing set

y_pred = linear_regression.predict(X_test)


# Evaluate the model's performance

mse = mean_squared_error(y_test, y_pred)

r2 = r2_score(y_test, y_pred)

print("Mean Squared Error:", mse)

print("R-squared Score:", r2)
```

**Output:**

```
warnings.warn(msg, category=FutureWarning)
Mean Squared Error: 24.291119474973463
R-squared Score: 0.6687594935356327
```

**9. Write a program to implement a decision tree classifier using scikit-learn and visualize the decision tree and understand its splits.**

```
import numpy as np

import pandas as pd

from sklearn.datasets import load_iris

from sklearn.model_selection import train_test_split

from sklearn.tree import DecisionTreeClassifier, plot_tree

import matplotlib.pyplot as plt

# Load Iris dataset

iris = load_iris()

X = iris.data

y = iris.target

# Split the dataset into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize Decision Tree classifier

decision_tree = DecisionTreeClassifier()

# Train the classifier

decision_tree.fit(X_train, y_train)
```

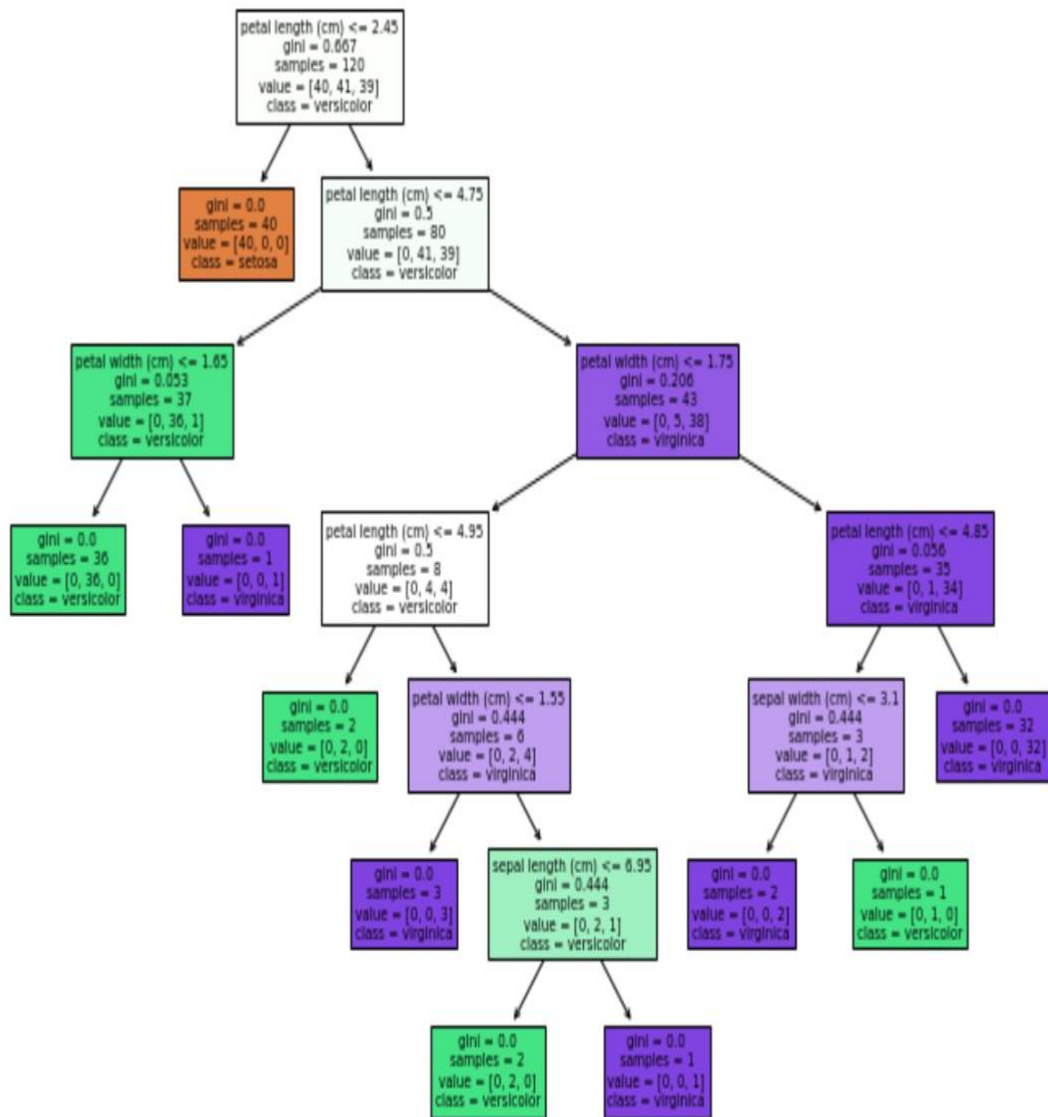
```
# Visualize the decision tree
```

```
plt.figure(figsize=(12, 8))
```

```
plot_tree(decision_tree, feature_names=iris.feature_names,  
class_names=iris.target_names, filled=True)
```

```
plt.show()
```

## Output:



## 10. Write a program to Implement K-Means clustering and Visualize clusters.

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import make_blobs
from sklearn.cluster import KMeans

# Generate sample data
X, y = make_blobs(n_samples=500, centers=4, cluster_std=0.8, random_state=42)

# Create a K-Means clusterer with 4 clusters
kmeans = KMeans(n_clusters=4, random_state=42)

# Fit the data
kmeans.fit(X)

# Get cluster labels
labels = kmeans.labels_

# Plot the data with cluster labels

plt.figure(figsize=(8, 6))
plt.scatter(X[:, 0], X[:, 1], c=labels, cmap='viridis')
plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_[:, 1], s=100,
            c='red', label='Centroids')
plt.title('K-Means Clustering')
plt.xlabel('X')
plt.ylabel('Y')
plt.legend()
plt.show()
```

**Output:**

